

Generalized linear regression with Pytho...

Saved to Dropbox • Dec 9, 2017 at 3:16 PM

# CODING TRICKS

MACHINE LEARNING  
AND CODING TRICKS

 [Leave a Comment](#)

 Updated on October 15, 2016

## Generalized linear regression with

# Python and scikit-learn library





One of the most used tools in machine learning, statistics and applied mathematics in general is

the regression tool. I say the regression, but there are lots of regression models and the one I will try to cover here is the well known generalized linear regression. The idea underneath this complex word is quite simple: the observation we try to fit is supposed to be a linear combination of input explanatory variables.

Mathematically, if the observation vector is called  $\hat{\mathbf{y}}$  this can be formulated by

$$\hat{\mathbf{y}} = \sum_k w_k \mathbf{x}_k$$

where the  $\mathbf{x}_k$  stands for the explanatory vector  $k$ , and the  $w_k$  are the weight of each explanatory variable. Furthermore, as we usually need for an intercept, by convention we set  $x_0 = 1$  so that the weight  $w_0$  represents the intercept.

# Linear Regression

As a first example, let's begin by a simple linear regression which consists of minimizing the sum of error squares. The error is defined as the difference between

the expected true value and the predicted value obtained by our model.

Mathematically, it expresses

$$\min_w \|X\mathbf{w} - \mathbf{y}\|_2^2$$

To illustrate this simple example, let's use the awesome library scikit-learn and especially the package


```
sklearn.linear_model
```

```
1  # General
2  import numpy as np
3  import pandas as pd
4  # Charting
5  from bokeh.plotting import figure
6  from bokeh.models import ColumnDataSource
7  from bokeh.io import output_notebook
8  # Regression
9  from sklearn.linear_model import LinearRegression
10
11
12 # For this example, I use output_notebook
13 output_notebook()
14
```

```

15 # Define data
16 x_1 = np.linspace(0,10,11)
17 x_0 = np.ones(len(x_1))
18 y = 4*x+1+np.random.normal(0,
19 data = pd.DataFrame(np.column_
20
21 # Define the model by explicit
22 reg = LinearRegression(fit_int
23
24 # Fit the data

```



The model we use here is quite simple, it is just a line. The model seems quite good with fitted coefficients of  $\beta_0$  and  $\beta_1$

$w_0 = -0.87998$  and  $\text{\textbackslash latex}$

$w_1 = 4.54914$ , but the error is not null

(mean squared error = 15.57 in my

example). In general in machine learning,

a way to reduce the residual error is to

change the model by a slightly more

complex one. In our case, we simply fitted

a polynomial of degree 1. What if we

increase the polynomial degree? For

example, let's say we increase the degree

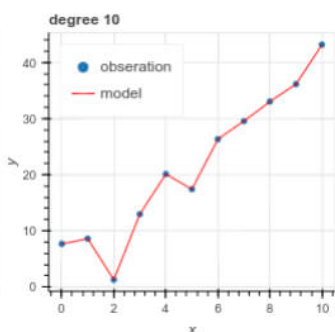




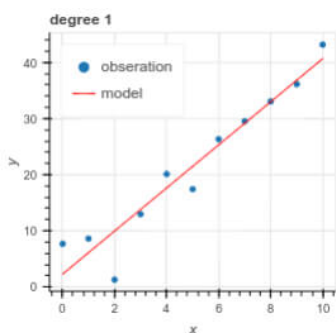
up to 12:

```
1 np.random.seed(10)
2 x_1=np.linspace(0,10,11)
3 x_0=np.ones(len(x_1))
4 y=4*x+1+np.random.normal(0, 5
5 data = pd.DataFrame(np.column_
6
7 predictors=['x_0', 'x_1']
8 predictors.extend(['x_{}'.format(i) for i in range(2,13):
9
10     data['x_{}'.format(i)] = (
11
12 col = ['rss', 'intercept'] + ['
13 ind = ['alpha_%.2g'%alpha for
14 coef_matrix_lasso = pd.DataFra
15
```

```

16 graphs = []
17 models = []
18 idx = 0
19 for degree in [1, 2, 4, 6, 8,
20               X = data[predictors[0:degree]]
21               reg = LinearRegression(fit
22               reg.fit(X, data['y'])
23               p=figure(width=300, height
24               p.circle(x,y,legend='obser

```





As we can see, the more we increase the polynomial degree of the model, the more we reduce the residual error. However, and it is particularly evident in this example, the reductions of the error is not necessarily a sign of a better model. Indeed, imagine we use a high degree polynomial as our model, the error tends to be null (and it is actually here as we have a polynomial degree equal to the number of observations). But if we add an extra observation, our model surely experiences a high residual error.

```
1 new_x = 15
2 data_x = np.array([new_x**i for i in range(15)])
3 new_y = 4*new_x+1+np.random.normal(0,1)
4 idx = 0
5 for degree in [1, 2, 4, 6, 8, 10, 12, 14]:
6     y_pred = models[idx].predict([new_x])
7     rss = (y_pred-new_y)**2
8     idx += 1
9     print "The residual error is: ", rss
```

and the result is

1	The residual error for $X=15$ is
2	The residual error for $X=15$ is
3	The residual error for $X=15$ is
4	The residual error for $X=15$ is
5	The residual error for $X=15$ is
6	The residual error for $X=15$ is

As we can see in this example, adding an observation at  $x=15$  leads to increasing error with the polynomial degree. This behavior is known as overfitting, i.e. the model fits very well the data but tends to poorly perform on new data. We say that it suffers of great variance.

To overcome this problem, we need to choose the right polynomial degree. We could for example split our dataset into two parts, a train set and a test set. Then, the best model would be the one with the least residual error on the test set.

However, there a clever method to limit the overfitting phenomenon: regularization.

# Regularization

Regularization consists of adding an penalty to a model, with the goal of preventing overfitting. It comes from the constataion that when the degree of the polynom increases (to take our first example), the weights of each monom also increases. Therefore, to overcome overfitting, we penalize monoms with high weight. The minimization function now becomes

$$\min_w \|X\mathbf{w} - \mathbf{y}\|_2^2 + \alpha \|\mathbf{w}\|^2$$

where  $\|\cdot\|$  is typically L1 or L2 norm, and  $\alpha$  is an hyper-parameter that can be tunable to adjust the penalty sensibility (0 means

no penalty, i.e. unregularized model). The two widely used regularization methods are L1 and L2 regularization, also called [lasso](#) and [ridge](#) regression.







# Lasso

To cite content of the scikit-learn library

---

Lasso is useful in some contexts due to its tendency to prefer solutions with fewer parameter values, effectively reducing the number of variables upon which the given solution is dependent. For this reason, the Lasso and its variants are fundamental to the field of compressed sensing. Under certain conditions, it can recover the exact set of non-zero weights.

```
1  # General
2  import numpy as np
3  import pandas as pd
4  # Charting
5  from bokeh.charts import show
6  from bokeh.models.glyphs import
7  from bokeh.plotting import fi
8  from bokeh.models import Colu
9  from bokeh.layouts import gri
10 from bokeh.io import output_r
11 # Regression
12 from sklearn.linear_model import
13 output_notebook()
14
15
16 def generate_plot(data, title=
17     source = ColumnDataSource(
18     p = Scatter(data, x='x', y
19     line = Line_glyph(x='x', y
20     p.add_glyph(source, line)
21     return p
22
23
24 def lasso_regression(data, pre
```




We clearly see the effect of regularization. When we increase the penalty, we strongly limit the weights of each coefficient, up to only keep the intercept. But how to choose the right parameter ? Well, here again, we need to look at the residual error computed on a set which is not the

training set. Such a method is known as validation. The principle is to split the data set into three parts, say 80% for the training set, 10% for validation and 10% for test. The model is trained on the training set, then the validation set is used to choose the hyper-parameter. Finally, the test set is used to estimate the true error of the model. However, on small data sets, this approach is not efficient enough as it limits the amount of data available for training. For such small data sets, we can apply the method of cross-validation. For that, we split the data set into two parts, one for the training and the other for the test. The training is then performed on all the training set but some  $k$  samples. So let's imagine that the training set is composed of  $N$  samples. We perform  $N/k$  regressions on the  $N-k$  samples and we compute the validation error on the  $k$  remaining samples. After all those regressions, the validation error is the mean error of all the validation errors.

In the scikit-learn library, there is a class that implement this approach and finds the optimal hyper-parameter: LassoCV. We then re-use the preceeding code sample omitting the `alpha` parameter to force the use of the `LassoCV` model :

```
1 graph, model = lasso_regressio
2 graph.title.text = 'alpha: {:.
3 show(graph);
```



According to the LassoCV model, the optimal hyper-parameter  $\alpha = 4.16e^{-4}$ .



# Ridge Regression

The ridge regression is quite similar to lasso, and differs only by the order of the norm used in the regularization term. In lasso, we used a norm of order 1 (L1) and in the ridge regression we use a norm of order 2 (L2). The behavior of this regularization technique is that all resulting weights of the models are still non null, but eventually with very small value so that their influence on the predicted value is quite low. On the opposite, lasso imposes sparsity of the model by eventually setting weights to null which make the model interpretation easier. The main advantage of the ridge regression is that it is indifferent to multiplicative factor, and tends to equals weights of highly-correlated variables whereas lasso will choose or the other.







The implementation of our example is really similar to the preceeding with lasso. All we have to do is replacing

Lasso	and	LassoCV	by
Ridge	and	RidgeCV	!


```
1  # General
2  import numpy as np
3  import pandas as pd
4  # Charting
5  from bokeh.charts import show
6  from bokeh.models.glyphs import
7  from bokeh.plotting import fi
8  from bokeh.models import Colu
9  from bokeh.layouts import gri
10 from bokeh.io import output_r
11 # Regression
12 from sklearn.linear_model import
13 output_notebook()
14
15
16 def generate_plot(data, title=
17     source = ColumnDataSource(
18     p = Scatter(data, x='x', y
19     line = Line_glyph(x='x', y
20     p.add_glyph(source, line)
21     return p
22
```

23

24

**def** ridge\_regression(data, pre

And with the use of cross-validation:



As we can see, the resulting model is a little different as with the lasso regularization. And if we now look at the weight coefficients:

Degree	1	2	3
L1	-0.569	-0.101	0
L2	-0.292e-01	-0.225e-01	-0.02

6	7	8	9
5.813e-04	0	0	0
3.389e-04	5.571e-05	6.083e-06	-6.7908

10	11	12	13
0	0	0	0
-2.72e-07	-9.600e-08	-2.353e-08	-4.4009

# Conclusion

In this article, we saw two regularization techniques and the importance to ALWAYS use regularization when we try to fit a model. We also saw that both techniques, although quite similar, give very different results. Indeed, whereas the

ridge technique includes all the explanatory variables, lasso results in a sparse model which is often easier to understand. However, lasso performs less well in the case of highly correlated variables as it tends to produce high sparsity in the result. And that's exactly what we saw in our example when we tried to fit a cosine with a polynomial. Each variable are highly correlated so that the resulting model has a lot of zero values as weight coefficients. Along with Ridge and Lasso, Elastic Net is another useful techniques which combines both L1 and L2 regularization. It allows for learning a sparse model while it also keep ridge properties.

## End note

We briefly evocate the subject, but as always in machine learning, we need several datasets to train the model, validate it and test it. We discussed that a little in this post and we quickly saw how to deal with small datasets with cross-

validation. This could be the subject of a next post.

---

### Like this:

Loading...

---

☰ Category: [Algorithm](#), [Machine Learning](#), [Python](#)    🏷️ Tags: [Lasso](#), [Machine Learning](#), [Python](#), [Regression](#), [Regularization](#), [Ridge](#)

---

← [Theano demo with Docker help](#)  
[Your first CNN made easy with Docker and Tensorflow](#) →

## Leave a Reply

# SEARCH

GO

## RECENT POSTS

---

[Your first CNN made easy with Docker and Tensorflow](#)

---

[Generalized linear regression with Python and scikit-learn library](#)

---

[Theano demo with Docker help](#)

---

[Slack Python client to send messages](#)

---

[Scrapping bank website to access account operations](#)

---

## RECENT COMMENTS

## ARCHIVES

---



November 2017

---

October 2016

---

September 2016

---

March 2016

---

October 2015

---

May 2015

---

April 2015

---

December 2014

---

November 2014

---

## CATEGORIES

---

.NET

---

Algorithm

---

C#

---

Database

---

Deep learning

---

[Docker](#)

---

[Google Service](#)

---

[JSON](#)

---

[Machine Learning](#)

---

[Matlab](#)

---

[MongoDB](#)

---

[Nodejs](#)

---

[Python](#)

---

[remote desktop](#)

---

[SCrapping](#)

---

[SQLite](#)

---

[Tips](#)

---

[WPF](#)

---

**META**

---

[Log in](#)

---

[Entries RSS](#)

---

[Comments RSS](#)

---

[WordPress.org](#)

---

## SHARE



## SEARCH

GO

Copyright © 2017 · All Rights Reserved · Coding  
Tricks

Swell Lite from Organic Themes · [RSS Feed](#) ·

[Log in](#)